


# End to End Data Project (ETL)

---



---

In this section, we'll build and execute an **ETL (Extract, Transform, Load)** workflow — a core data engineering process used to move data from its source to a destination where it can be analyzed and visualized.

- **Extract** – Retrieve raw data from various sources such as APIs, databases, or files.
- **Transform** – Clean, format, and structure the data to make it consistent and meaningful.
- **Load** – Store the processed data into a target system, such as a database or data warehouse.

Our **first step** will be to **extract the data** from the source. We'll connect to the dataset, verify access and structure, and ensure it's ready for transformation in the next phase.

# Part 1: EXTRACT

---

## Summary

- Initial environment setup in PowerShell
  - Creating Kaggle API credentials
  - Install tools and launch Jupyter lab from PowerShell
  - Create Jupyter notebook and verify Kaggle CLI access and data set
- 

## Configure Environment in Powershell

### 1. Launch PowerShell & Navigate to the Project

- Open Windows PowerShell (Start Menu → type PowerShell → Enter).
- Choose or create a folder for your project, for example:

```
cd C:\Users\<YourUser>\Projects  
mkdir end-to-end-retail-analytics  
cd end-to-end-retail-analytics
```

### 2. Create and Activate a Virtual Environment

- Create a virtual environment:

```
python -m venv .venv
```

- Activate it:

```
.\.venv\Scripts\Activate
```

- Confirm the environment is active: your prompt should now start with (.venv) .

### 3. Install Required Python Packages

- Install the core tools needed for extraction and exploration:

```
pip install kaggle jupyterlab pandas
```

- You can also install any additional libraries later as needed.
-

# Set up Kaggle API Credentials

## Kaggle README Instructions

### API credentials

To use the Kaggle API, sign up for a Kaggle account at <https://www.kaggle.com>. Then go to the 'Account' tab of your user profile (<https://www.kaggle.com/<username>/account>) and select 'Create API Token'. This will trigger the download of `kaggle.json`, a file containing your API credentials. Place this file in the location appropriate for your operating system:

- Linux: `$XDG_CONFIG_HOME/kaggle/kaggle.json` (defaults to `~/.config/kaggle/kaggle.json`). The path `~/.kaggle/kaggle.json` which was used by older versions of the tool is also still supported.
- Windows: `C:\Users\<Windows-username>\.kaggle\kaggle.json` - you can check the exact location, sans drive, with `echo %HOMEPATH%`.
- Other: `~/kaggle/kaggle.json`

You can define a shell environment variable `KAGGLE_CONFIG_DIR` to change this location to `$KAGGLE_CONFIG_DIR/kaggle.json` (on Windows it will be `%KAGGLE_CONFIG_DIR%\kaggle.json`).

For your security, ensure that other users of your computer do not have read access to your credentials. On Unix-based systems you can do this with the following command:

```
chmod 600 ~/.config/kaggle/kaggle.json
```

You can also choose to export your Kaggle username and token to the environment:

```
export KAGGLE_USERNAME=datadinosaur
export KAGGLE_KEY=xxxxxxxxxxxxxx
```



In addition, you can export any other configuration value that normally would be in the `kaggle.json` in the format 'KAGGLE\_-' (note uppercase). For example, if the file had the variable "proxy" you would export `KAGGLE_PROXY` and it would be discovered by the client.

## 4. Configure Kaggle API (One-Time Setup)

### 4.1 Get Your Kaggle API Token

- Go to Kaggle in your browser and sign in.
- Click your profile icon → Settings.
- Scroll to API → click Create New API Token.
- A file named `kaggle.json` will download.

### 4.2 Place `kaggle.json` in the Correct Location (Windows)

- In File Explorer, go to:

`C:\Users\<YourUser>\`

- If it doesn't exist, create a folder:

`.kaggle`

- Move `kaggle.json` into:

`C:\Users\<YourUser>\.kaggle\`

- (Optional but recommended) Right-click `kaggle.json` → Properties → check Hidden → OK  
This keeps the key out of casual view.

## 5. Verify Kaggle CLI Configuration

Back in your activated PowerShell session:

```
kaggle datasets list -s retail
```

- If you see a list of datasets, your API is configured correctly.
- If you get an error about credentials or kaggle not found:
  - Make sure kaggle.json is in C:\Users\<YourUser>\.kaggle\
  - Make sure you're in the .venv and kaggle is installed (pip install kaggle)

## 6. Download the Dataset with Kaggle API

Inside your project folder:

- Create a folder to store raw data:

```
mkdir data_raw
```

- Download and unzip the dataset (example uses a retail orders dataset):

```
kaggle datasets download -d ankitbansal06/retail-orders -p  
data_raw --unzip
```

- -d specifies the dataset.
- -p data\_raw saves files into the data\_raw folder.
- --unzip extracts the zip automatically.

At this point, the Extract from Kaggle → local environment is complete.

## 7. Launch JupyterLab from the Project Environment

From the same folder, with .venv still active:

```
jupyter lab
```

- This will open JupyterLab in your browser.
- Make sure the URL points to your project directory.

If prompted to choose a kernel, select the one associated with .venv.

## 8. Connect and Inspect the Data in Jupyter

In JupyterLab:

- Create a new Python 3 notebook.
- Run:

```
import pandas as pd
```

```
file_path = "data_raw/Retail Orders.csv" # update to the actual  
filename  
df = pd.read_csv(file_path)  
  
df.head()  
df.info()
```

This confirms:

- The Kaggle API connection worked
  - The dataset is successfully extracted and readable
  - The notebook is running inside the correct environment
-

## Part 2: Transform


---

### Description:

Once your data is successfully loaded, the **next phase** is **Transform**, where you'll:

- Clean missing or inconsistent values,
  - Format data types,
  - Derive new columns or metrics, and
  - Prepare the data for analysis or loading into a database.
- 

### Set Up Jupyter Environment



#### Step 1: Create your new notebook

You're in the **Launcher** tab, which is perfect.

Under "**Notebook**", click on:

Kaggle Analytics

That will open a new tab named something like `untitled.ipynb`.

This ensures that your notebook runs inside the **virtual environment** (`venv`) where you installed your packages (pandas, matplotlib, kaggle, etc.).

### Step 2: Rename the notebook

In the file browser on the left, right-click the file `Untitled.ipynb` → click Rename → type:

```
explore.ipynb
```

This keeps things organized and descriptive.



### Step 3: Verify the kernel

Look at the top-right corner of the notebook (beside the circle icon).

You should see:

Kaggle Analytics

If it says anything else (like "Python 3 (ipykernel)"), click on it and change it to Kaggle Analytics.

That's your environment with all your installed libraries.

### Step 4: Run your first cell to confirm setup

In the first cell, paste and run this:

```
python
import pandas as pd
import matplotlib.pyplot as plt
import sqlite3

print("Setup successful!")
print("Pandas version:", pd.__version__)
```

Then hit Shift + Enter to execute the cell.

If it prints the lines correctly with no errors, your environment is fully working and connected to the right Python.

Confirmation Example below.

The screenshot shows the Jupyter Notebook interface. On the left, the file browser displays a directory structure with files: `data`, `venv`, `explore.ipynb` (selected), and `requirements.txt`. In the main area, a code cell [1] contains the following Python code:

```
import pandas as pd
import matplotlib.pyplot as plt
import sqlite3

print("Setup successful!")
print("Pandas version:", pd.__version__)
```

The cell output shows the results of the execution:

```
[1]: Setup successful!
Pandas version: 2.3.3
```

## Step 5: (Optional Check) Confirm Kaggle CLI access

In a new cell, run:

```
python
```

Copy code

```
!kaggle --version
```

You should see something like:

```
nginx
```

Copy code

```
Kaggle API 1.6.x
```

That confirms your Kaggle token and CLI tool are installed correctly in this environment.

```
[2]: !kaggle --version
```

Kaggle API 1.7.4.5

## Step 6: Load your dataset

Now that everything's ready, we'll start analysis.

You already downloaded your dataset earlier (the retail orders CSV).

Let's verify and load it.

```
python
```

Copy

```
df = pd.read_csv("data/Retail_Orders.csv")
df.head()
```

If you see a table appear, that means the dataset loaded perfectly and we can move on to exploration, cleaning, and SQL integration.

```
[5]: df = pd.read_csv("data/orders.csv")
df.head()
```

|   | Order Id | Order Date | Ship Mode      | Segment   | Country       | City            | State      | Postal Code | Region | Category        | Sub Category | Product Id      | cost price | List Price | Quantity | Discount Percent |
|---|----------|------------|----------------|-----------|---------------|-----------------|------------|-------------|--------|-----------------|--------------|-----------------|------------|------------|----------|------------------|
| 0 | 1        | 2023-03-01 | Second Class   | Consumer  | United States | Henderson       | Kentucky   | 42420       | South  | Furniture       | Bookcases    | FUR-BO-10001798 | 240        | 260        | 2        | 2                |
| 1 | 2        | 2023-08-15 | Second Class   | Consumer  | United States | Henderson       | Kentucky   | 42420       | South  | Furniture       | Chairs       | FUR-CH-10000454 | 600        | 730        | 3        | 3                |
| 2 | 3        | 2023-01-10 | Second Class   | Corporate | United States | Los Angeles     | California | 90036       | West   | Office Supplies | Labels       | OFF-LA-10000240 | 10         | 10         | 2        | 5                |
| 3 | 4        | 2022-06-18 | Standard Class | Consumer  | United States | Fort Lauderdale | Florida    | 33311       | South  | Furniture       | Tables       | FUR-TA-10000577 | 780        | 960        | 5        | 2                |
| 4 | 5        | 2022-07-13 | Standard Class | Consumer  | United States | Fort Lauderdale | Florida    | 33311       | South  | Office Supplies | Storage      | OFF-ST-10000760 | 20         | 20         | 2        | 5                |

## Part 3: Clean

### Description:

After extracting the dataset from Kaggle, the next step was to clean and prepare it for analysis. The raw data contained missing values, inconsistent formatting, and unnecessary columns. Using Pandas in JupyterLab, I performed a series of data cleaning tasks to ensure accuracy and consistency across all records.

### Key Steps:

1. Loaded the dataset into a Pandas DataFrame for inspection using `df.info()` and `df.describe()`.

2. Handled missing values by filling or removing null entries depending on the column's significance.
3. Standardized data types for date, numeric, and categorical fields to ensure compatibility for SQL import and calculations.
4. Removed duplicates and irrelevant columns that did not contribute to the analysis.
5. Renamed columns for readability and uniformity (e.g., converting spaces to underscores).
6. Created new derived columns such as total revenue (quantity \* unit\_price) and profit margin to support later analysis.
7. Validated the dataset by re-checking row counts, null values, and unique identifiers before loading it into SQL Server.

This process produced a clean, structured dataset ready for the Transform and Load phases of the pipeline.

---

### Data Cleaning Process for this project data:

#### Handle Null Values

```
[17]: #Handle Null Values
df = pd.read_csv("../data/orders.csv",na_values=['Not Available','unknown'])
df['Ship Mode'].unique()

[17]: array(['Second Class', 'Standard Class', nan, 'First Class', 'Same Day'],
       dtype=object)
```

#### Format Columns

```
[21]: #rename columns, make them lower case, and replace space with underscore
df.rename(columns={'Order Id':'order_id', 'City':'city'})
df.columns=df.columns.str.lower()
df.columns=df.columns.str.replace(' ','_')
df.head(5)
```

| Order<br>Id | Order<br>Date | Ship<br>Mode | Segment | Country | to | order_id | order_date | ship_mode | segment | country |
|-------------|---------------|--------------|---------|---------|----|----------|------------|-----------|---------|---------|
|-------------|---------------|--------------|---------|---------|----|----------|------------|-----------|---------|---------|

#### Derive new columns

```
[23]: #derive new columns discount, sale price, and profit
df['discount']=df['list_price']*df['discount_percent']*.01
df['sale_price']= df['list_price']-df['discount']
df['profit']=df['sale_price']-df['cost_price']
df
```

| cost_price | list_price | quantity | discount_percent | discount | sale_price | profit |       |
|------------|------------|----------|------------------|----------|------------|--------|-------|
| 240        | 260        | 2        |                  | 2        | 5.2        | 254.8  | 14.8  |
| 600        | 730        | 3        |                  | 3        | 21.9       | 708.1  | 108.1 |
| 10         | 10         | 2        |                  | 5        | 0.5        | 9.5    | -0.5  |

#### Check Data Types

```
[24]: #check data types  
df.dtypes
```

```
[24]: order_id          int64  
order_date        object  
ship_mode         object  
segment           object  
country            object  
city               object  
state              object  
postal_code       int64  
region             object  
category           object  
sub_category      object  
product_id        object  
cost_price         int64  
list_price         int64  
quantity           int64  
discount_percent   int64  
discount           float64  
sale_price         float64  
profit             float64  
dtype: object
```

## Convert Order Date Data Type

```
[25]: #convert order date from object data type to datetime  
df['order_date']=pd.to_datetime(df['order_date'],format="%Y-%m-%d")
```

```
[26]: df.dtypes
```

```
[26]: order_id          int64  
order_date        datetime64[ns]
```

## Drop Unnecessary Columns

```
#drop cost price list price and discount percent columns  
df.drop(columns=['list_price','cost_price','discount_percent'],inplace=True)
```

```
df.head(5)
```

|  | quantity | discount | sale_price | profit |
|--|----------|----------|------------|--------|
|  | 2        | 5.2      | 254.8      | 14.8   |
|  | 3        | 21.9     | 708.1      | 108.1  |

---

# Part 3: Load

## Description:

### System Requirements:

- SQL Server Express
- SQL Server management Studio (SSMS)

## Load Process for this project:

Set up column data types in SSMS to minimize storage volume.

The screenshot shows the SQL Query Editor in SSMS. The code pane contains the following SQL script:

```
1 CREATE TABLE df_orders (
2     order_id INT PRIMARY KEY,
3     order_date DATE,
4     ship_mode VARCHAR(20),
5     segment VARCHAR(20),
6     country VARCHAR(20),
7     city VARCHAR(20),
8     state VARCHAR(20),
9     postal_code VARCHAR(20),
10    region VARCHAR(20),
11    category VARCHAR(20),
12    sub_category VARCHAR(20),
13    product_id VARCHAR(50),
14    quantity INT,
15    discount DECIMAL(7, 2),
16    sale_price DECIMAL(7, 2),
17    profit DECIMAL(7, 2)
18 );
19
```

The status bar at the bottom indicates "Commands completed successfully." and the completion time: 2025-11-01T15:41:18.3288336-05:00.

Upload data into SSMS database using “append” to ...

```
1: import sqlalchemy as sal
2: from sqlalchemy import text

3: server      = r"localhost\SQLEXPRESS"
4: driver      = "ODBC Driver 17 for SQL Server"
5: target_db   = "e2e_data_project"

6: # 1) Connect to master with AUTOCOMMIT so CREATE DATABASE is allowed
7: engine_master = sal.create_engine(
8:     f"mssql+pyodbc:///{server}/master?driver={driver.replace(' ', '+')}&trusted_connection=yes",
9:     isolation_level="AUTOCOMMIT", # <-- key: no transaction
10)

11 with engine_master.connect() as conn:
12     # Check existence first (single statement)
13     exists = conn.execute(text("SELECT DB_ID(:name)"), {"name": target_db}).scalar()
14     if exists is None:
15         # CREATE DATABASE must be its own statement with autocommit
16         conn.exec_driver_sql(f"CREATE DATABASE [{target_db}]")

17 # 2) Connect to the target DB (normal transactional behavior)
18 engine = sal.create_engine(
19     f"mssql+pyodbc:///{server}/{target_db}?driver={driver.replace(' ', '+')}&trusted_connection=yes"
20 )

21 # 3) Load your dataframe
22 with engine.begin() as conn:
23     df.to_sql("df_orders", con=conn, index=False, if_exists="append")

24 print("Database ensured and data uploaded.")

25 Database ensured and data uploaded.
```

Verify data loaded and data constraints are correct.

The screenshot shows the SSMS interface with the following details:

- Object Explorer:** On the left, it shows the database structure for "localhost\SQLEXPRESS (SQL Server 16.0.1150 - LAPTOP)". Under the "e2e\_data\_project" database, the "Tables" node is expanded, showing the "dbo.df\_orders" table. The "Columns" node under "df\_orders" is also expanded, listing columns: order\_id, order\_date, ship\_mode, segment, country, city, state, postal\_code, region, category, sub\_category, product\_id, quantity, discount, sale\_price, and profit.
- SQL Query Window:** On the right, titled "SQLQuery2.s... \edmys (59)\*", contains the query: "select \* from df\_orders".
- Results Window:** Below the query window, the results are displayed in a grid format. The columns are order\_id, order\_date, ship\_mode, and segment. The data consists of 17 rows:

| order_id | order_date | ship_mode      | segment     |
|----------|------------|----------------|-------------|
| 1        | 2023-03-01 | Second Class   | Consumer    |
| 2        | 2023-08-15 | Second Class   | Consumer    |
| 3        | 2023-01-10 | Second Class   | Corporate   |
| 4        | 2022-06-18 | Standard Class | Consumer    |
| 5        | 2022-07-13 | Standard Class | Consumer    |
| 6        | 2022-03-13 | NULL           | Consumer    |
| 7        | 2022-12-28 | Standard Class | Consumer    |
| 8        | 2022-01-25 | Standard Class | Consumer    |
| 9        | 2023-03-23 | NULL           | Consumer    |
| 10       | 2023-05-16 | Standard Class | Consumer    |
| 11       | 2023-03-31 | NULL           | Consumer    |
| 12       | 2023-12-25 | NULL           | Consumer    |
| 13       | 2022-02-11 | Standard Class | Consumer    |
| 14       | 2023-07-18 | Standard Class | Consumer    |
| 15       | 2023-11-09 | NULL           | Home Office |
| 16       | 2022-06-18 | Standard Class | Home Office |
| 17       | 2022-02-04 | Standard Class | Consumer    |

A message at the bottom of the results window says "Query executed successfully."